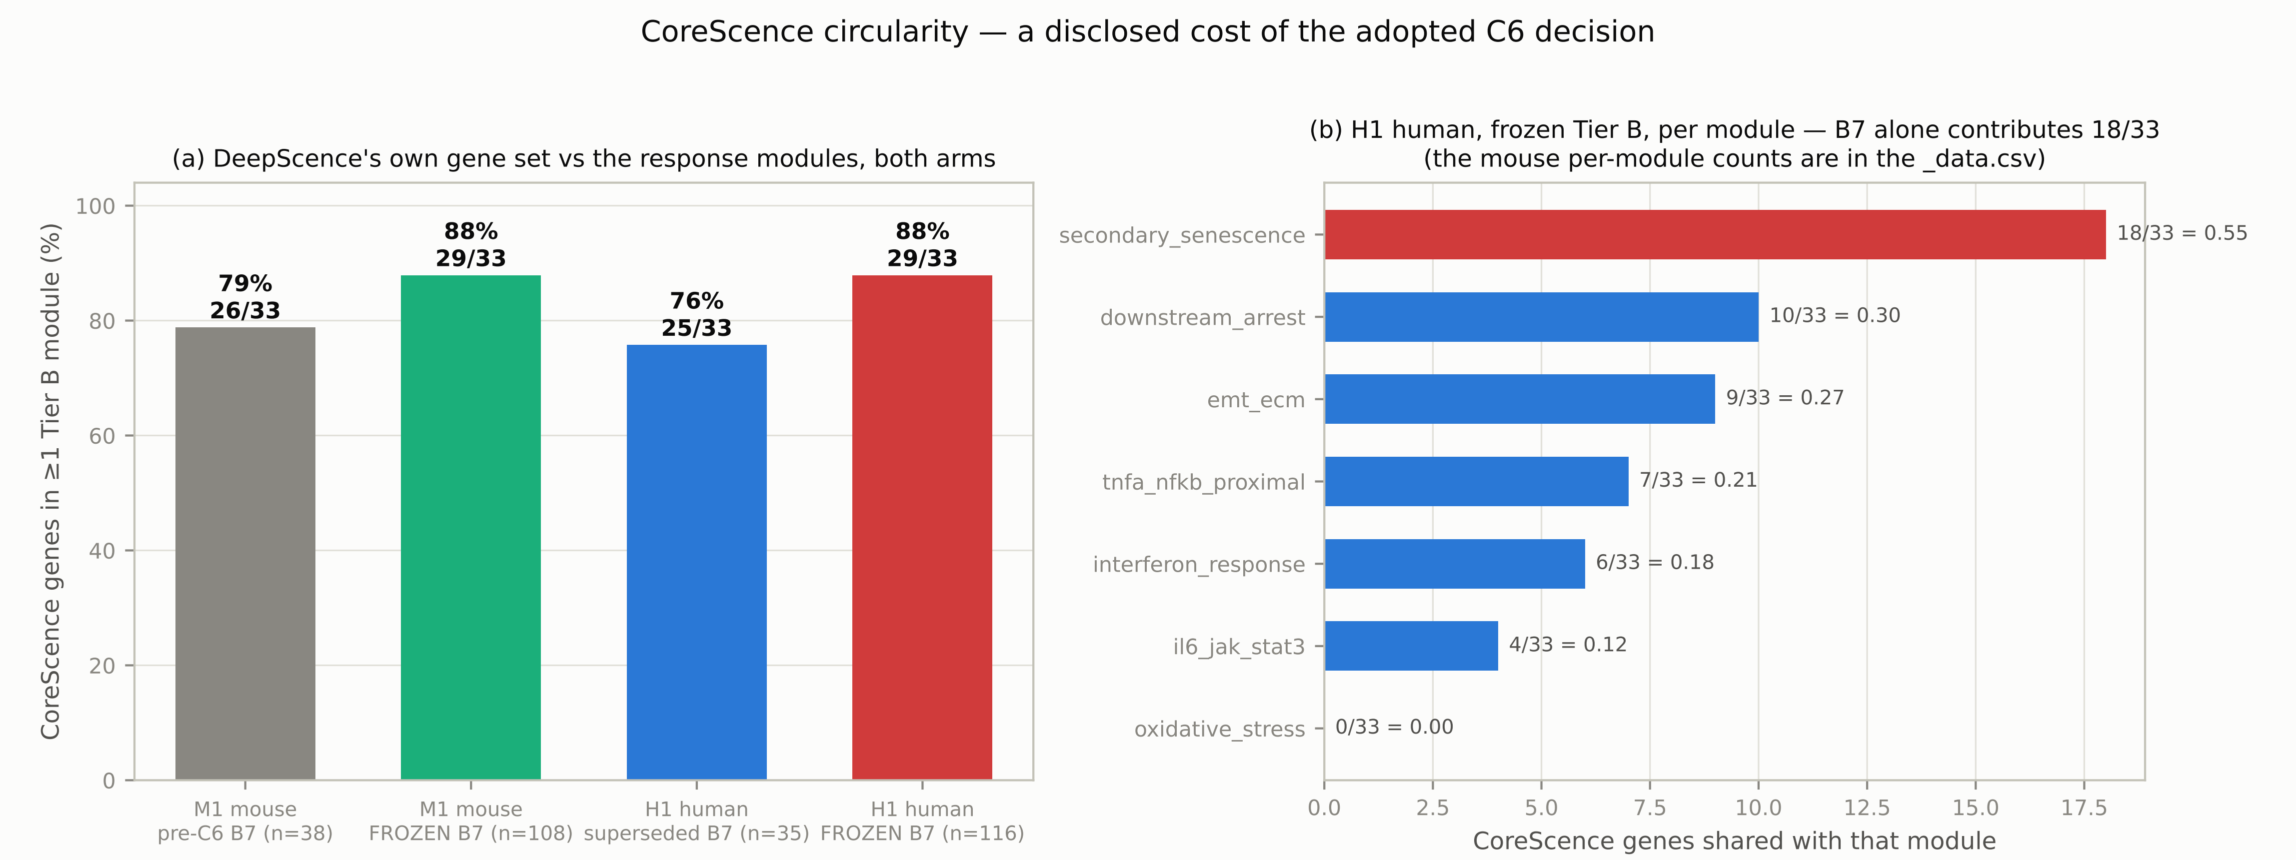

DeepScence is the sender caller for Job B, so overlap between its gene set and the readouts is circularity by construction. Re-sourcing B7 from SenMayo ∪ Reactome SASP fixed the sender/readout collision but raised this figure on BOTH arms: mouse 79% → 88%, human 76% → 88%. CITE 88% — it is the configuration that will be run. The mouse bars are the 33 CoreScence genes reachable on the 5,097-gene mouse panel (31 through the pinned 1:1 MGI map plus CDKN2B and CXCL1, which the map has no row for); the same convention code/run_phase3_n8.py uses, and the denominator in the committed results/phase3/n8_disjointness_*.csv. An earlier version of this panel showed a typed-in 24/35 = 69% for the mouse arm; that denominator is reproducible under no convention and understated the pre-C6 mouse circularity by ~10 points. CoreScence is human, so on H1 it runs natively and none of the human numbers is ortholog-mapping loss. Strip-and-refit is in the frozen run order.